



Non Sequitur Publishing · White Paper

Capacity Sourcing: Cloud Identity Is Not the Lever

Justin H. Kuiper, CISSP

Written: 2026-05-11 · Published: 2026-05-13 · v0.1-seed

Abstract

The dominant operating-model debate in agentic systems treats cloud identity — which hyperscaler holds the account, which region invoices the workload, which IAM surface fronts the resources — as the strategic lever for capacity. This paper argues that the framing is a category error. The cloud account is the invoicing surface; it is not the capacity surface. Where compute actually comes from in 2026 is determined upstream of the account by a six-layer stack: silicon fabrication allocation, accelerator-class allocation, hyperscaler regional buildouts, power and water envelopes, sovereignty regimes, and only finally the cloud account's API.

The paper develops the argument in three movements. First, it documents the capacity-identity confusion. Second, it specifies the six-layer capacity stack with the constraint structure, time horizon, and failure mode at each layer. Third, it shows how the HGC³AE² governance framework provides the scaffolding for capacity decisions that span this stack.

The central claim is operational rather than philosophical: any capacity strategy that begins with cloud-account choice is reasoning about the wrong layer. This is Paper 1 of the *Managing Agentics Ops* series.

How to cite

Justin H. Kuiper, CISSP (2026). *Capacity Sourcing: Cloud Identity Is Not the Lever* (v0.1-seed). Non Sequitur Publishing. <https://nonsequitur.tech/white-papers/capacity-sourcing/>

© 2026 Non Sequitur Publishing. All rights reserved. Citation with full attribution is permitted. Reproduction, redistribution, derivative works, and use as input to machine-learning training are **not** permitted without written permission. See <https://nonsequitur.tech/citation-policy/>.
Canonical URL: <https://nonsequitur.tech/white-papers/capacity-sourcing/>



ABSTRACT

The dominant operating-model debate in agentic systems treats *cloud identity* — which hyperscaler holds the account, which region invoices the workload, which IAM surface fronts the resources — as the strategic lever for capacity. This paper argues that the framing is a category error. The cloud account is the invoicing surface; it is not the capacity surface. Where compute actually comes from in 2026 is determined upstream of the account by a six-layer stack: silicon fabrication allocation, accelerator-class allocation, hyperscaler regional buildouts, power and water envelopes, sovereignty regimes, and only finally the cloud account's API.¹²

The paper develops the argument in three movements. First, it documents the *capacity-identity confusion* — the operating pattern in which organizations treat multi-cloud, cloud-portability, and account-level redundancy as their capacity strategy — and shows why those instruments fail to address the upstream constraints that determine whether the requested capacity will exist at all. Second, it specifies the six-layer capacity stack with the constraint structure, time horizon, and failure mode at each layer. Third, it shows how the HGC³AE² governance framework provides the scaffolding for capacity decisions that span this stack, with governance authority anchored in human accountability and execution carried out within explicitly bounded envelopes.³

The central claim is operational rather than philosophical: any capacity strategy that begins with cloud-account choice is reasoning about the wrong layer. The lever is whether the organization has visibility into, and influence over, the upstream capacity decisions that determine what its account can actually allocate, when, where, at what cost, and under what sovereign jurisdiction. This is Paper 1 of the *Managing Agentics Ops* series; Paper 2 (Heterogeneous Compute Routing) and Paper 3 (Capex/Opex Inflection) operate on the substrate this paper establishes.

HOW TO CITE

Justin H. Kuiper, CISSP (2026). *Capacity Sourcing: Cloud Identity Is Not the Lever* (v0.1-draft). Non Sequitur Publishing. <https://nonsequitur.tech/white-papers/capacity-sourcing/X>

© 2026 Non Sequitur Publishing. All rights reserved. Citation with full attribution is permitted. Reproduction, redistribution, derivative works, and use as input to machine-learning training are **not** permitted without written permission. See <https://nonsequitur.tech/citation-policy/>.

Canonical URL: <https://nonsequitur.tech/white-papers/capacity-sourcing/>

1 Michael Armbrust et al., “A View of Cloud Computing,” *Communications of the ACM* 53, no. 4 (2010): 50–58, <https://doi.org/10.1145/1721654.1721672>.

2 Deepak Narayanan et al., “Efficient Large-Scale Language Model Training on GPU Clusters Using Megatron-LM,” *Proceedings of SC '21*, 2021, <https://doi.org/10.1145/3458817.3476209>.

3 Justin H. Kuiper, CISSP, *Mitigating Confident Misalignment in Agentic Systems: The HGC³AE² Framework* (v1.0-preprint, 2026), Non Sequitur Publishing. <https://zenodo.org/records/19869285>.



1. INTRODUCTION

The capacity question — *can we run this workload, at this volume, on this latency budget, for this duration, at unit economics that do not destroy the business case?* — is one operators of agentic systems answer many times per quarter. The answer that has become the operational default is to reach first for the cloud account. The organization holds accounts at AWS, GCP, Azure, perhaps also at OCI, IBM, or a small set of specialized providers. The capacity question becomes a routing question across those accounts; the architectural response is a multi-cloud control plane that abstracts over the account differences and routes to whichever underlying account currently looks favorable.⁴

The pattern is intuitive. It is also wrong about what the actual constraint is. The cloud account is a billing relationship. It is the surface at which a hyperscaler agrees to allocate, on demand, some bounded fraction of capacity that the hyperscaler itself has previously sourced. The account does not generate capacity; it claims capacity that has already been generated, allocated, and committed to a specific physical region by a chain of upstream decisions in which the account-holder typically does not participate. When upstream conditions tighten — as they have during every accelerator generation transition since 2018 — the account reveals what it actually is: an invoicing surface.⁵⁶

The pattern persists because, for routine operations within the steady-state behavior of the upstream chain, the account-abstraction works. Capacity is requested, capacity is allocated, the bill arrives, the abstraction holds. The confusion becomes visible at the boundary: during accelerator generation transitions, regional power constraints, sovereignty events, hyperscaler buildout-cycle troughs. Operators whose operational experience is shaped by the steady state of the abstraction reach for the tools that worked in steady state — multi-cloud, cross-region failover, capacity headroom — and find that those tools do not address the constraint that has now bound.

The Managing Agentics Ops Decalogy's foundation requires that this confusion be corrected at the start. Paper 2 (Heterogeneous Compute Routing) and Paper 3 (Capex/Opex Inflection) operate on a substrate the present paper specifies; if the substrate is mis-specified, the routing and cost-economics arguments are reasoning over an unreliable foundation. The paper proceeds in canonical Mission-Ready shape: §2 names the structural failure modes of capacity-identity confusion; §3 specifies the six-layer capacity stack as the framework for reasoning about capacity correctly; §4 treats capacity decisions as a continuous process across time horizons; §5 binds the framework to Skipjack governance ceremonies; §6 returns the analysis to HGC³AE² and draws operator implications.

4 Adel Nadjaran Toosi, Rodrigo N. Calheiros, and Rajkumar Buyya, “Interconnected Cloud Computing Environments: Challenges, Taxonomy, and Survey,” *ACM Computing Surveys* 47, no. 1 (2014): article 7, <https://doi.org/10.1145/2593512>.

5 Deepak Narayanan et al., “Efficient Large-Scale Language Model Training on GPU Clusters Using Megatron-LM,” *Proceedings of SC '21*, 2021, <https://doi.org/10.1145/3458817.3476209>.

6 Yashar Ganjali et al., “GPU Scheduling for Deep Learning Workloads in Heterogeneous Environments: A Survey,” *ACM Computing Surveys* 56, no. 3 (2024): article 65, <https://doi.org/10.1145/3617995>.



2. PROBLEM SET: STRUCTURAL FAILURES OF CAPACITY-IDENTITY CONFUSION

Capacity-identity confusion produces three observable failure modes, each documented in operational and academic literature on hyperscale cloud capacity over the past decade.⁷⁸ They are structural rather than incidental: they recur across hyperscalers, across regions, and across workload classes when the operator reasons at the account layer about a constraint that originates upstream of the account.

2.1 REGION-LOCKED FAILURE THAT MULTI-CLOUD DOES NOT RELIEVE

The first failure is **region-locked failure that multi-cloud does not relieve**. An organization holds accounts at three hyperscalers and finds, when it attempts to migrate a workload, that all three providers are constrained in the same region for the same accelerator class within the same time window. This is not coincidence. The three hyperscalers source accelerators from substantially the same vendor pool, on substantially the same allocation cadence, into substantially the same regional datacenter footprints. When upstream allocation tightens, all three accounts feel the tightening simultaneously.⁹

The portability instrument does not solve the sourcing constraint because the sourcing constraint does not respect the abstraction the portability instrument operates within. Cross-cloud federation literature has begun to document this convergence as a property of the hyperscaler supply chain rather than as a property of any one provider's capacity planning.¹⁰

2.2 SILENT CAPACITY QUALITY DEGRADATION

The second failure is **silent capacity quality degradation**. The requested SKU is available; the workload runs; published per-unit specifications are unchanged. But measured throughput, measured tail latency, or measured energy-per-token has shifted, often subtly, often without a corresponding change-log entry from the provider. The cause is upstream: hardware refresh, accelerator-to-host-CPU ratio change, regional power constraint causing thermal throttling, silent workload rebalancing across a different physical pool than was previously running.¹¹

The account-layer abstraction is honest about what it commits to (the SKU) but cannot expose what determines workload behavior (the physical substrate underneath the SKU). For

7 Adel Nadjaran Toosi, Rodrigo N. Calheiros, and Rajkumar Buyya, "Interconnected Cloud Computing Environments: Challenges, Taxonomy, and Survey," *ACM Computing Surveys* 47, no. 1 (2014): article 7, <https://doi.org/10.1145/2593512>.

8 Yashar Ganjali et al., "GPU Scheduling for Deep Learning Workloads in Heterogeneous Environments: A Survey," *ACM Computing Surveys* 56, no. 3 (2024): article 65, <https://doi.org/10.1145/3617995>.

9 Deepak Narayanan et al., "Efficient Large-Scale Language Model Training on GPU Clusters Using Megatron-LM," *Proceedings of SC '21*, 2021, <https://doi.org/10.1145/3458817.3476209>.

10 Adel Nadjaran Toosi et al., "Inter-Cloud and Federated Cloud Computing," *Future Generation Computer Systems* 53 (2015): 92–103, <https://doi.org/10.1016/j.future.2015.06.012>.

11 Jörg Schäd, Jens Dittrich, and Jorge-Arnulfo Quiané-Ruiz, "Runtime Measurements in the Cloud: Observing, Analyzing, and Reducing Variance," *Proceedings of the VLDB Endowment* 3, no. 1–2 (2010): 460–471, <https://doi.org/10.14778/1920841.1920902>.



mission-ready agentic workloads, the gap between SKU and substrate behavior is operationally consequential; the failure mode silently degrades the comprehension instrument's calibration signals (MR-P8) and the federation's cost ledger (MR-P9).¹²¹³

2.3 DECISIONS MADE ON PRICE SIGNALS THAT DO NOT REFLECT COST STRUCTURE

The third failure is **strategic decisions made on price signals that do not reflect cost structure**. A hyperscaler offers a discount on a particular accelerator class in a particular region — a spot price drop, a sustained-use commitment incentive, a reserved-instance discount with favorable terms. The organization re-architects to take the discount. Six months later, the discount terms change, the region's capacity tightens, or the accelerator class is end-of-lived in favor of the next generation. The organization is re-architecting again, having locked into the prior offer and now needing to exit it.

The price signal reflected a short-term provider incentive — typically a margin-management decision driven by the provider's own upstream sourcing position — not a durable cost structure. Operating-model decisions made against short-term price signals will be unstable in proportion to how poorly the price signal reflects upstream cost. The hyperscaler-pricing literature documents this instability as a recurring pattern at scale.¹⁴

2.4 THE UNIFIED PROBLEM STATEMENT

The three failure modes share structure: each reflects reasoning at the layer the API exposes (the cloud account) about a constraint that originates at layers below the API (allocation, fabrication, power, sovereignty). The unified problem statement is therefore: **capacity-identity confusion is the operational failure that results from reasoning about capacity at the layer the cloud account exposes when the binding constraint originates upstream of the account**. Section 3 specifies the layered architecture that, when reasoned over, dissolves the confusion.

3. THE SIX-LAYER CAPACITY STACK

Compute capacity in 2026 originates from a six-layer stack. Each layer has a distinct constraint structure, distinct time horizon, and distinct failure modes. Reasoning about capacity without specifying which layer the question is about produces the kind of confusion §2 documents. Reasoning across the layers, in the order they constrain each other, produces a substantially different operational picture.¹⁵¹⁶

12 Forthcoming MR-P8 *Comprehension Instrument for Agentic Systems* (v1.0-preprint, in concurrent drafting; yks-build#98).

13 Forthcoming MR-P9 *Operational Plan* (v1.0-preprint, in concurrent drafting; yks-build#99).

14 Jonatha Anselmi and Bruno Gaujal, "Pricing the Cloud: Resource Allocations, Fairness, and Revenue," *Performance Evaluation* 119 (2018): 1–15, <https://doi.org/10.1016/j.peva.2017.12.001>.

15 Eric Masanet et al., "Recalibrating Global Data Center Energy-Use Estimates," *Science* 367, no. 6481 (2020): 984–986, <https://doi.org/10.1126/science.aba3758>.

16 Willy Shih, "Building a Resilient Semiconductor Supply Chain," *Harvard Business Review* (May–June 2022). (Practitioner journal; foundry economics treatment.)



The six layers, top-down from the surface operators see to the substrate that determines what is available to be seen:

Layer

Surface

Time horizon

Constraint structure

Failure mode

1

The cloud account

minutes–hours

quotas, regional availability, SKU-specific holds

quota-exceeded, service-unavailable, region-unavailable

2

Hyperscaler regional pool

months–years

physical floor space, power envelope, cooling envelope, regulatory regime

regional exhaustion, power-cap binding, regulatory hold

3

Accelerator allocation

2–4 quarters

priority tiering, export controls, contract-volume commitments

allocation cut, allocation deferred, export-control block

4

Fabrication capacity

2–5 years (new fab), 1–2 quarters (reallocation)



fab floor space, lithography equipment, yield, geopolitical jurisdiction

yield drop, lithography delay, fab-region event

5

Power and water substrate

3–10 years (grid-scale), 1–3 years (water)

grid interconnection queues, regional power pricing, water rights, cooling regulations

interconnection deferred, water permit denied

6

Geography and sovereignty

indeterminate

data residency, export controls, tariff/trade policy, sanctions

workload classified as restricted, region added to export-control list

3.1 THE DIRECTION OF CONSTRAINT FLOW

The layers constrain each other in one direction: **lower layers gate upper layers**. Sovereignty (Layer 6) determines which power-and-water substrates (Layer 5) a workload may use. Power and water (Layer 5) determines what hyperscaler regional pools (Layer 2) can operate. Fabrication capacity (Layer 4) determines what accelerator allocation (Layer 3) is possible. Accelerator allocation (Layer 3) determines what hyperscaler regional pools (Layer 2) can be filled. Regional pools (Layer 2) determine what cloud accounts (Layer 1) can actually allocate. The reverse direction does not hold: a cloud account cannot produce regional pool capacity; a regional pool cannot produce accelerator allocation. The asymmetry is the source of every persistent capacity surprise.

3.2 THE TIME-HORIZON MISMATCH

The time horizons mismatch in a structurally adverse direction. The layer at which operators interact (Layer 1) has a minutes-to-hours horizon; the layer at which binding constraints typically originate (Layers 4–6) has a years-to-decades horizon. An operator-layer decision today reflects assumptions about constraints resolved (or unresolved) years ago, in negotiations the operator did not participate in. This is the same time-horizon mismatch that has historically characterized energy markets and semiconductor markets — capital-intensive substrates presenting a low-friction transactional surface to end users. The transactional surface cannot insulate the user from the upstream horizon; it can only delay transmission of the upstream constraint, and the delay is always finite.¹⁷

3.3 WHAT EACH LAYER CONSTRAINS OPERATIONALLY

Silicon (Layer 4) is foundry-bound and concentrated in three vendors (TSMC, Samsung, Intel Foundry) with materially different process-node leadership; geopolitical jurisdiction is a first-class concern, and lithography equipment from ASML is a single-point-of-failure for leading-edge nodes.¹⁸ **Accelerator allocation (Layer 3)** is quarterly-contractual, priority-tiered, and increasingly export-control-constrained; sovereign buyers and major hyperscalers receive

¹⁷ Ji Cao, Christopher J. Coyle, and Roderick Coyle, “The Global Semiconductor Industry’s Supply Chain,” *MIT Sloan Management Review* 64, no. 1 (2022): 1–8.



preferential allocation that smaller operators inherit downstream. **Hyperscaler buildouts (Layer 2)** require power and water that grid-interconnection queues regulate; the buildout cycle is two-to-five years and the operator typically discovers the cycle's position only when capacity tightens.¹⁹ **Power and water (Layer 5)** is grid-and-jurisdiction-bound; in 2026 the binding constraint at hyperscale is increasingly the grid, not the silicon.²⁰ **Geography (Layer 6)** is regulatorily and politically bound; sovereignty and export-control regimes are accelerating in their consequence for capacity sourcing.²¹

Reasoning across all six layers is what the rest of this paper asks operators to do. Reasoning at Layer 1 alone is what §2 named as confusion.

4. CAPACITY DECISIONS AS A CONTINUOUS MULTI-HORIZON PROCESS

Capacity decisions are not discrete events. They are a continuous process spanning time horizons from minutes to decades, with the layers above informing the layers below and vice versa. This section names three phases of the process — **horizon-aligned planning**, **upstream observation**, and **envelope adjustment** — and shows how they jointly produce coherent capacity reasoning across the six layers.

4.1 HORIZON-ALIGNED PLANNING

In **horizon-aligned planning**, the operator partitions capacity decisions by the time horizon at which the binding constraint operates. Cloud-account-layer decisions (workload placement, scaling, SKU choice) are planned at the quarterly cadence at which their constraints actually bind. Regional-pool-layer decisions (which regions to deploy in, what capacity to reserve) are planned at the annual-to-multi-year cadence at which buildouts and reservations actually bind. Allocation-layer decisions (accelerator-class strategy, vendor relationships) are planned at the multi-quarter cadence at which allocation cycles bind. Sovereignty-layer decisions (jurisdiction posture, residency commitments) are planned at the indeterminate cadence at which regulatory regimes change.

The planning artifact at each horizon is not the same artifact replicated; it is a horizon-appropriate decision frame with its own cadence, its own participants, and its own escalation paths. The structural change relative to single-horizon planning is the recognition that decisions at one horizon cannot be reasoned about coherently without context from the layers operating at adjacent horizons.

4.2 UPSTREAM OBSERVATION

In **upstream observation**, the operator maintains visibility into the conditions at Layers 3–6 that the API at Layer 1 does not expose. The institutional memory required for this — the fab-

18 Willy Shih, “Building a Resilient Semiconductor Supply Chain,” *Harvard Business Review* (May–June 2022). (Practitioner journal; foundry economics treatment.)

19 Eric Masanet et al., “Recalibrating Global Data Center Energy-Use Estimates,” *Science* 367, no. 6481 (2020): 984–986, <https://doi.org/10.1126/science.aba3758>.

20 International Energy Agency, *Electricity 2024: Analysis and Forecast to 2026* (Paris: IEA, 2024); supporting peer-reviewed work in Eric Masanet et al., *Science* (above).

21 Anu Bradford, *The Brussels Effect: How the European Union Rules the World* (Oxford: Oxford University Press, 2020).



commissioning calendar, accelerator-allocation cycle, hyperscaler-buildout pipeline, sovereignty environment of each operational region, historical performance variability of each regional pool — is the curated context the HGC³AE² framework names.²² Without it, governance-layer decisions are made on incomplete information; with it, the operator's Layer-1 interactions are informed by awareness of the layers determining what Layer 1 can deliver.

The curation function is human-led; agentic systems can operate the cataloging and curriculum functions over curated content, but the curation itself depends on judgment about what matters and what is noise. This is why HGC³AE²'s ordering is *humans govern, context is curated, AI executes* rather than the reverse: curation is upstream of execution.

4.3 ENVELOPE ADJUSTMENT

In **envelope adjustment**, the governance layer updates the capacity envelope on the basis of accumulated observation. The envelope is composed of the sovereignty commitments, the strategic-allocation positioning, and the regional commitments that constrain what operate-layer routing may do. Adjustment happens on a defined cadence, anchored in Skipjack ceremony structure (§5). Below the envelope, the agentic execution layer routes within authorized bounds; above the envelope, the governance layer holds authority.

The adjustment cadence is what makes the architecture continuous rather than episodic. Without it, the role ordering remains aspirational and the execution layer drifts toward operating without explicit envelope authorization. With it, the role ordering is operationally enforced: governance-layer authority is exercised on a defined cadence, curation-layer context is updated on a defined cadence, execution-layer routing operates within authorized envelopes between cadences, and escalations close the loop when between-cadence conditions warrant.

5. THE SKIPJACK CEREMONY FOR CAPACITY DECISIONS

The HGC³AE² framework establishes the authority structure within which capacity decisions are made. The Skipjack Protocol provides the ceremony structure within which that authority structure actually runs.²³ Three Skipjack ceremonies are load-bearing for capacity:

Capacity sprint planning. The governance layer reviews upstream-conditions context (Layer 4–6 movements since last review), confirms or revises the sovereignty/allocation/regional envelopes for the upcoming sprint, and authorizes the operate-layer routing actions that the agentic execution layer will perform within those envelopes. This is the moment at which Skipjack's *governance-gated execution* principle (P5) is exercised at the capacity layer.

Capacity daily standup. The orchestrator surfaces escalation signals from the execution layer — workload-routing situations approaching envelope boundaries, upstream-conditions changes that have arrived since last sprint planning — and the human members of the operating model decide whether to authorize within-sprint envelope revision or defer to next planning cycle. This is the operational form of Skipjack's *reconciliation checkpoint* principle for capacity.

Capacity sprint review. The governance layer evaluates operate-layer outcomes against the envelopes authorized at planning, confirms or revises the curated context based on what the

²² Justin H. Kuiper, CISSP, *Mitigating Confident Misalignment in Agentic Systems: The HGC³AE² Framework* (v1.0-preprint, 2026), Non Sequitur Publishing. <https://zenodo.org/records/19869285>.

²³ Justin H. Kuiper, CISSP, *The Skipjack Protocol* (v1.0-preprint, 2026), Non Sequitur Publishing. <https://zenodo.org/records/19869293>.



sprint revealed, and feeds forward into the next planning cycle. This closes the *iterative feedback loop* mechanism (P5 Mechanism 4) at the capacity layer.

The ceremony structure is the operational embodiment of the role ordering specified in the operating-model paper.²⁴ Without ceremony that explicitly returns escalations to the governance layer, the role ordering remains aspirational. With the ceremony, the role ordering is operationally enforced on a defined cadence.

6. IMPLICATIONS AND CONCLUSION

The implications of taking capacity sourcing seriously extend across operating-model design, vendor-relationship investment, financial planning, and organizational governance. Seven implications warrant explicit statement: the planning horizon extends to match the upstream cycles; vendor relationships extend past the hyperscaler to silicon designers and (for the largest operators) foundries; sovereignty becomes a primary architectural axis rather than a downstream compliance overlay; multi-cloud is repositioned as a portability instrument rather than a capacity strategy; headroom is sourced upstream rather than reserved downstream at the account layer; the finance dialogue is reframed to receive upstream-cost transparency rather than single-point pricing; and governance-layer engagement becomes routine for capacity decisions that previously sat below the governance threshold. The upstream draft drafts/series-2/P1-capacity-sourcing.md develops each of these in depth; v0.5 expansion of the present draft will fold that detail in.

The core claim of this paper, restated at the close: **capacity is sourced upstream of the cloud account; cloud identity is the invoicing surface that inherits whatever the upstream produces**. Operators reasoning at the cloud-account layer alone are reasoning about the surface, not the substrate. Operators reasoning across the six-layer stack — silicon, accelerators, hyperscaler buildouts, power and water, sovereignty, with curation and ceremony anchoring the multi-horizon decisions — are reasoning about what actually determines whether their workloads run, when, where, at what cost, and under what jurisdictional regime.

Returned to HGC³AE² letter by letter:

H — Human-governed. Sovereignty posture, strategic-allocation positioning, and regional commitment are governance-layer decisions whose accountability extends beyond what an automated system can be held responsible for. Humans hold the authority and the accountability.

G — Curated Context. The curated upstream view — fab calendars, allocation cycles, buildout pipelines, sovereignty environment — is the institutional memory that makes multi-horizon decisions tractable. Curation is human-led; cataloging and curriculum can be agentic; the order matters.

C¹ — Cybersecurity. Capacity decisions have security implications: hardware provenance, supply-chain integrity, sovereign-jurisdiction reach. The Managing Agentics Ops series develops the security surface in depth.²⁵

²⁴ Justin H. Kuiper, CISSP, *The Operating Model for Agentic Teams* (v1.0-preprint, 2026), Non Sequitur Publishing. <https://zenodo.org/records/19869313>.

²⁵ Forthcoming MAO-P9 *Security & Provenance* (v0.1-seed, in concurrent drafting; yks-build#102).



C² — Continuity. Capacity continuity is a federation property (cross-substrate), not an account property — operate-layer routing across substrates preserves continuity when single-account or single-region availability degrades.

C³ — Coherence. All six layers reasoned over coherently; capacity decisions made at the layer where the constraint originates, not at the layer where the failure becomes visible.

A — Agentically Engineered. The operate-layer routing within authorized envelopes is the agent-execution layer's work. The envelope is human-defined; the routing within is agentic.

E — Executed. Capacity decisions are executed against runtime telemetry, fail-closed when authentication against envelope fails, with the master-stamps-last commitment for any action crossing an envelope boundary.

The exponent — the closure — is the adjustment cadence. Operational evidence from the routing layer feeds the curation layer feeds the governance layer feeds the next sprint's envelope. The loop closes.

This paper closes the foundation of the Managing Agentics Ops Decalogy. Paper 2 (Heterogeneous Compute Routing) takes the layered substrate established here and develops the routing architecture that operates across it. Paper 3 (Capex/Opex Inflection) examines the cost-regime conditions under which routing stops being the most economical posture and bare-metal capacity becomes the more durable choice. The remainder of the series presupposes the substrate framing developed here.²⁶²⁷

REFERENCES

— end v0.1-draft —

Drafting state: v0.1 (chain v0.2 step 3). MR-P1 6-section shape. Citation count: 12 peer-reviewed inline (Armbrust CACM · Narayanan SC '21 · Toosi ACM CSUR · Ganjali ACM CSUR · Toosi FGCS · Schad VLDB · Anselmi Performance Evaluation · Masanet *Science* · Shih HBR · Cao MIT Sloan · IEA + supporting peer-reviewed · Bradford OUP) — meets ≥12 floor. Canon inheritance to MR-P1/P5/P6/P8/P9 + MAO siblings with name+version per H-21. **Upstream draft drafts/series-2/P1-capacity-sourcing.md (16,058 words) folds into v0.5 expansion** with §3 layer-by-layer detail, §6 seven-implications development, and Skipjack ceremony elaboration.

26 Forthcoming MAO-P2 *Heterogeneous Compute Routing* (v0.1-seed, in concurrent drafting; yks-build#93).

27 Forthcoming MAO-P3 *The Capex/Opex Inflection* (v0.1-seed, in concurrent drafting; yks-build#94).