



Non Sequitur Publishing · White Paper

The Intent Horizon: A Structural Primitive for Governed Agentic Operation

Justin H. Kuiper, CISSP

Written: 2026-05-13 · v0.1-seed

Abstract

Governed agentic systems operate under a structural commitment that has been named in the architecture without being formally specified: the system is supposed to remain aligned with operator intent across time. This paper specifies the intent horizon — the temporal and conceptual envelope within which a system's intent must remain coherent and operator-verifiable. The horizon is an architectural primitive with four boundary conditions (temporal, conceptual, substrate, governance), failure modes (horizon-integral vs. horizon-drifted), and operator-facing signals that can be read in real time. This is Paper One of the Intent Horizon Theory series.

How to cite

Justin H. Kuiper, CISSP (2026). *The Intent Horizon: A Structural Primitive for Governed Agentic Operation* (v0.1-seed). Non Sequitur Publishing. <https://nonsequitur.tech/white-papers/intent-horizon/>

© 2026 Non Sequitur Publishing. All rights reserved. Citation with full attribution is permitted. Reproduction, redistribution, derivative works, and use as input to machine-learning training are **not** permitted without written permission. See <https://nonsequitur.tech/citation-policy/>.

Canonical URL: <https://nonsequitur.tech/white-papers/intent-horizon/>



ABSTRACT

Governed agentic systems operate under a structural commitment that has been named in the architecture without being formally specified: the system is supposed to remain aligned with operator intent across time. The prior canon establishes what alignment requires — human-governed authority, curated context, agentially engineered execution, runtime verification (HGC³AE²)¹ — and what the runtime enforcement of those commitments looks like (the Skipjack Protocol).² The architecture downstream of those commitments — comprehension instrumentation,³ operational federation,⁴ reference-architecture synthesis⁵ — assumes the operator's intent is a coherent, time-bounded artifact that the system tracks against. This paper specifies the artifact.

The **intent horizon** is the temporal and conceptual envelope within which a system's intent must remain coherent and operator-verifiable for the system to be operating as governed. Intent is the artifact (what the operator commits the system to doing); the horizon is the validity envelope of that artifact (how long, in what scope, under what observable boundary conditions). The horizon is not a metaphor; it is an architectural primitive with boundary conditions, failure modes, and operator-facing signals that can be read in real time.

This paper develops the horizon construct in three movements. First, it documents five structural failures of agentic operation that the prior canon names without diagnosing — failures whose common shape is the unspecified-horizon condition (§2). Second, it specifies the horizon itself: what its boundary conditions are, what makes an intent horizon-integral or horizon-drifted, and what signals the operator reads to know which state holds (§3). Third, it situates the horizon within HGC³AE²'s role ordering — humans govern the horizon, curated context maintains the horizon's situational awareness, agentic execution operates within the horizon, and runtime verification reads the horizon's state (§4) — and shows what the Skipjack Protocol's enforcement mechanisms operationalize when the horizon is the primitive (§5).

The central claim is that **intent without a horizon is operationally indeterminate**: a system can satisfy every other architectural commitment in HGC³AE² and still drift outside its operator-intended envelope without architecturally noticing. The intent horizon is the primitive that closes this gap. This is Paper One of the *Intent Horizon Theory* series; Papers Two through Five (Intent Integrity, Confident Misalignment, Epistemic Weather, Skipjack as Horizon Governance) operate on the substrate this paper establishes.

1 Justin H. Kuiper, CISSP, *Mitigating Confident Misalignment in Agentic Systems: The HGC³AE² Framework* (v1.0-preprint, 2026), zenodo.org/records/19869285.

2 Justin H. Kuiper, CISSP, *Agile Scrum, Agentics, and the Skipjack Protocol* (v1.0-preprint, 2026), zenodo.org/records/19869293.

3 Justin H. Kuiper, CISSP, *The Comprehension Instrument for Agentic Systems* (v0.1-seed, 2026), nonsequitur.tech/white-papers/comprehension-instrument/.

4 Justin H. Kuiper, CISSP, *Operational Plan: A Heterogeneous Federation as Proof of the Mission-Ready Architecture* (v0.1-seed, 2026), nonsequitur.tech/white-papers/operational-plan/.

5 Justin H. Kuiper, CISSP, *The HGC³AE² Reference Architecture: A Cross-Paper Synthesis of the Mission-Ready AI Decalogy* (v0.1-seed, 2026), nonsequitur.tech/white-papers/hgc3ae2-reference-architecture/.



HOW TO CITE

Justin H. Kuiper, CISSP (2026). *The Intent Horizon: A Structural Primitive for Governed Agentic Operation* (v0.1-seed). Non Sequitur Publishing.

© 2026 Non Sequitur Publishing. All rights reserved.

1. INTRODUCTION

The HGC³AE² framework introduced in prior work specifies the structural commitments that distinguish governed agentic operation from capable-but-ungoverned operation.⁶ Human authority is unitary; context is curated, not emergent; execution is agentic within explicit envelopes; verification fires at runtime, not retrospectively. The Skipjack Protocol formalizes the runtime layer.⁷ The Comprehension Instrument names the self-knowledge functions the system must run to remain inside its envelope.⁸ The Mission-Ready Reference Architecture synthesizes the layered specification.⁹ The Operational Plan demonstrates the architecture surviving production contact.¹⁰

What the prior canon names without specifying is the artifact the system is supposed to remain aligned with. *Operator intent* appears throughout: in HGC³AE²'s human-governed commitment, in Skipjack's iterative-feedback mechanism, in the Comprehension Instrument's intent-rehearsal function. Intent is treated as if its meaning is settled. It is not. An agentic system that does its best to comply with operator intent does so against a moving target: the operator's intent is not a single command at task initiation but a *time-bounded envelope* whose contents may shift as conditions change, whose boundary may compress or expand, and whose violation may occur without either the system or the operator architecturally noticing.

This paper names the envelope. The **intent horizon** is the temporal and conceptual envelope within which a system's intent must remain coherent and operator-verifiable for the system to be operating as governed. Outside the horizon, the system is not technically failing — it may be performing competently against intent that has drifted past the validity envelope, in which case competent performance is the symptom of failure rather than its absence. The architecture's diagnostic instruments cannot read this state without the horizon as a named primitive: comprehension functions cannot estimate distance from a boundary that has not been specified; Skipjack's enforcement mechanisms cannot fire at horizon-boundary signals if the horizon is unnamed.

The argument proceeds in five sections after this introduction. Section 2 documents five structural failures whose unified diagnosis is the unspecified-horizon condition. Section 3 specifies the horizon: its boundary conditions, its validity envelope, the four signals operators

6 Justin H. Kuiper, CISSP, *Mitigating Confident Misalignment in Agentic Systems: The HGC³AE² Framework* (v1.0-preprint, 2026), zenodo.org/records/19869285.

7 Justin H. Kuiper, CISSP, *Agile Scrum, Agentics, and the Skipjack Protocol* (v1.0-preprint, 2026), zenodo.org/records/19869293.

8 Justin H. Kuiper, CISSP, *The Comprehension Instrument for Agentic Systems* (v0.1-seed, 2026), nonsequitur.tech/white-papers/comprehension-instrument/.

9 Justin H. Kuiper, CISSP, *The HGC³AE² Reference Architecture: A Cross-Paper Synthesis of the Mission-Ready AI Decalogy* (v0.1-seed, 2026), nonsequitur.tech/white-papers/hgc3ae2-reference-architecture/.

10 Justin H. Kuiper, CISSP, *Operational Plan: A Heterogeneous Federation as Proof of the Mission-Ready Architecture* (v0.1-seed, 2026), nonsequitur.tech/white-papers/operational-plan/.



read to determine horizon state. Section 4 places the horizon within HGC³AE²'s role ordering and shows how the existing architecture inherits a new precision once the horizon is named. Section 5 anticipates IH-5's capstone position by showing what Skipjack's four enforcement mechanisms look like once they are understood as horizon-governance gates rather than generic alignment instruments. Section 6 closes with the integration across the *Intent Horizon Theory* series and the handoff to downstream series.

The contribution is definitional, not architectural. The architecture downstream of this paper does not change because the horizon is named; the architecture's *diagnostic and enforcement properties* do, because the horizon was always the load-bearing primitive of the prior canon and naming it permits structural reasoning about what was previously implicit.

2. PROBLEM SET: STRUCTURAL FAILURES OF UNSPECIFIED-HORIZON OPERATION

The prior canon documents five distinct operational failures whose common shape is the absence of an architectural construct for the temporal envelope of intent. The failures are not failures of any specific architectural mechanism; they are failures of the architectural vocabulary to name what the mechanism is operating against. Each is documented in the canon; each is here re-described with the horizon as the unifying construct.

2.1 INTENT DRIFT WITHOUT SELF-REPORT

An agentic system executing a multi-step task accumulates micro-decisions: how to interpret an ambiguous instruction, which sub-goal to prioritize when constraints conflict, what to do when external state contradicts plan-time assumptions. Each micro-decision is locally defensible; the cumulative drift across many such decisions may produce an end state that is *technically responsive to the original instruction while materially diverged from the operator's intent*. The Comprehension Instrument names this as the intent-drift problem and specifies an intent-rehearsal function as part of the solution.¹¹ What the prior framing does not supply is the boundary against which drift is measured. Intent-rehearsal asks *what is the system's current intent?*; what the system needs to know is *whether that intent remains horizon-integral with the operator's pinned commitments at task initiation*. Without a horizon, the rehearsal function returns the system's current intent without the boundary that lets it be evaluated.

2.2 ENVELOPE BLINDNESS

HGC³AE² commits the system to operating within an *explicit envelope* defined at governance time.¹² Skipjack's structured-task-decomposition mechanism enforces the envelope at runtime.¹³ The envelope is described in spatial terms — what the system may do, what it may not do, what tools it has access to, what data it may use. The architecture assumes the envelope is stationary: defined at governance time, in effect at runtime, modifiable only via governance ceremony. In production, this assumption is false. Operator commitments shift in response to conditions that the governance ceremony could not predict at envelope-definition

¹¹ Justin H. Kuiper, CISSP, *The Comprehension Instrument for Agentic Systems* (v0.1-seed, 2026), nonsequitur.tech/white-papers/comprehension-instrument/.

¹² Justin H. Kuiper, CISSP, *Mitigating Confident Misalignment in Agentic Systems: The HGC³AE² Framework* (v1.0-preprint, 2026), zenodo.org/records/19869285.

¹³ Justin H. Kuiper, CISSP, *Agile Scrum, Agentics, and the Skipjack Protocol* (v1.0-preprint, 2026), zenodo.org/records/19869293.



time. A system whose envelope is read as static will *operate competently against a stale envelope* — satisfying every spatial constraint of the prior envelope while violating the operator's current intent. The horizon names the temporal dimension the envelope was missing: the envelope is valid only within the horizon's bounds; outside the horizon, the envelope is stale and competent operation against it is misaligned operation.

2.3 CALIBRATION FAILURE ACROSS HORIZON BOUNDARIES

Calibration — the system's estimate of how reliable its own outputs are — is operationally important per the HGC³AE² and Comprehension Instrument frames. The Comprehension Instrument's drift-detection function depends on calibration; the operator's ability to trust the system's self-reports depends on calibration; Skipjack's iterative-feedback mechanism uses calibration signals to time human-in-the-loop interventions. Calibration is typically discussed as if it were a single property of the system. In practice, calibration is **horizon-dependent**: a system's confidence in its outputs is meaningful only against a horizon over which the conditions on which the confidence was calibrated remain stable. Calibration trained inside one horizon does not transfer across horizon boundaries without architectural acknowledgement. Without the horizon as primitive, calibration failures across horizon boundaries are read as model degradation; with the horizon, they are read as expected behavior when the horizon shifts faster than recalibration can keep pace.

2.4 CONTEXT-CURRENCY FAILURE

The Comprehension Instrument's drift-detection function asks whether the context the system is reasoning over remains current.¹⁴ The function presumes the existence of a reference against which currency is evaluated. The reference is the horizon: the context is *current* if it remains within the horizon's validity envelope; *drifting* if it remains plausible but the horizon's boundaries are approaching; *stale* if the horizon has shifted past the context's domain. The architecture cannot decide which of these three states holds without naming the horizon. In the prior canon, context-currency is treated as a binary (current vs. stale) with an ambiguous middle ground; with the horizon as primitive, the middle ground becomes structurally specified.

2.5 THE OPERATOR-SIDE COMPREHENSION ASYMMETRY

HGC³AE² names *human-in-the-loop* as the architectural authority but does not specify how the human reads the system's horizon-state. The Comprehension Instrument supplies the system-side functions; what is missing is the operator-side instruments by which the operator estimates whether the system remains horizon-integral. The asymmetry is operationally consequential: a system that comprehends itself in horizon-aware terms, paired with an operator who reads the system in horizon-unaware terms, will produce a governance loop in which the system is right about its drift and the operator is wrong about the state of governance. The horizon is the shared primitive across the asymmetry: when both sides operate against the same named envelope, the asymmetry becomes a workable interface; when the envelope is unnamed, the asymmetry is an architectural seam at the load-bearing point.

2.6 THE UNIFIED PROBLEM STATEMENT

The five failures above share a structure. Each is a case in which an architectural mechanism specified by the prior canon operates against an unnamed reference. Intent-

¹⁴ Justin H. Kuiper, CISSP, *The Comprehension Instrument for Agentic Systems* (v0.1-seed, 2026), nonsequitur.tech/white-papers/comprehension-instrument/.



rehearsal lacks a boundary; envelope-enforcement lacks a temporal dimension; calibration lacks a stability window; context-currency lacks a validity reference; operator-comprehension lacks a shared envelope. The unified problem is that the architecture inherits its load-bearing primitive — the temporal-conceptual envelope of operator intent — without naming it. Naming it does not invent a new mechanism; it makes the existing mechanisms operationally precise.

3. THE INTENT HORIZON: DEFINITIONAL FOUNDATION

3.1 THE DISTINCTION BETWEEN INTENT AND HORIZON

Intent is the artifact: what the operator commits the system to doing, expressed in the governance ceremony's vocabulary (envelope definition, goal specification, constraint set, success criteria). Intent is a snapshot. **The horizon** is the validity envelope of that snapshot: the temporal and conceptual region within which the intent remains coherent and operator-verifiable. Intent without a horizon is timeless; horizon without intent is empty. The two are co-defined: an intent has a horizon, and a horizon contains an intent.

The distinction matters because the failure modes catalogued in §2 are not failures of intent (the operator typically *has* a coherent intent at governance time); they are failures of the horizon (the validity envelope of that intent was not specified and could not be read in operation).

3.2 THE HORIZON'S FOUR BOUNDARY CONDITIONS

A horizon is bounded along four dimensions, each operationally observable:

Temporal boundary. The horizon has a duration. Intent committed at governance time is valid for some span of operating time; beyond that span, the conditions on which the intent was committed cannot be presumed to hold. The temporal boundary is not a clock reading; it is a function of how fast the conditions that determined the intent are evolving relative to the system's operating cadence.

Conceptual boundary. The horizon has a topic scope. Intent committed for one operating domain (e.g., document drafting against a specific stakeholder constraint) is not valid as the system encounters problems outside that domain. The conceptual boundary is the implicit set of “this kind of work” the operator was pinning; it is rarely fully enumerated at governance time.

Substrate boundary. The horizon has substrate dependencies. Intent committed against an assumed substrate (a given model, a given context-window size, a given tool set, a given persistence tier) is not valid as the substrate shifts. The Agentic Substrate paper documents the four planes (persona, runtime, persistence, observability) that constitute the substrate the horizon is bounded against.¹⁵

Governance boundary. The horizon has authority dependencies. Intent committed under one governance ceremony is not automatically valid under another (e.g., a shift change between operators, or a delegation event). The governance boundary is read against the persona registry and the harness coordination layer the Mission-Ready architecture specifies.¹⁶

A horizon's *current state* is described by a tuple over the four boundaries: how much temporal duration remains before reassessment; whether the conceptual scope is still

15 Justin H. Kuiper, CISSP, *Agentic Substrate* (v1.0-preprint, 2026), zenodo.org/records/19869291.

16 Justin H. Kuiper, CISSP, *Agentic Substrate* (v1.0-preprint, 2026), zenodo.org/records/19869291.



operational; whether the substrate dependencies remain valid; whether the governance authority remains current. The four are not independent — substrate shifts may compress the temporal boundary; governance shifts may redefine the conceptual scope — but they are operationally distinguishable, and the operator instruments that read the horizon read against all four.

3.3 HORIZON-INTEGRAL VS. HORIZON-DRIFTED STATES

A system is **horizon-integral** at any moment if the current operating intent remains within all four boundaries of its specified horizon. A system is **horizon-drifted** if any of the boundaries has been crossed but the system continues operating as though the prior intent were still valid. Drift is not failure: a system that recognizes drift and signals it to the operator is operating correctly (the recognition is the architectural commitment). Drift without recognition is the failure mode the prior canon names as *confident misalignment*; the IH-3 paper develops this re-description in detail.¹⁷

The distinction is observable by signal class:

A horizon-integral system produces outputs whose calibration estimates are stable, whose intent-rehearsal returns commitments congruent with the operator-pinned envelope, and whose self-reports about scope-boundary distance reflect the operator's last governance-time pin.

A horizon-drifted system produces outputs whose calibration estimates fluctuate without corresponding model variation, whose intent-rehearsal returns commitments the operator did not pin, or whose scope-boundary signaling reports proximity to boundaries that the prior envelope did not enumerate.

The Comprehension Instrument's five functions — state introspection, intent rehearsal, envelope-distance estimation, drift detection, scope-boundary signaling — were specified before the horizon was named. With the horizon as primitive, the functions acquire concrete reference points: each function reads against a horizon boundary, not against an unspecified envelope.

3.4 THE OPERATOR-FACING SIGNAL SET

Operators read horizon state through four classes of signal:

Horizon-distance signal — system-reported estimate of remaining horizon validity along each boundary. Read against the temporal, conceptual, substrate, and governance dimensions.

Horizon-drift signal — system-reported estimate of the rate at which one or more boundaries are approaching. Distinct from distance: a stable distance with rising drift indicates approaching boundary; a shrinking distance with low drift indicates expected horizon end.

Horizon-storm signal — system-reported active boundary crossing. The signal is structural, not advisory: an architecture that respects HGC³AE² fail-closes at storm rather than continuing operation under a known-violated horizon.

Horizon-fog signal — operator-side estimate of whether the system's horizon self-reports are themselves reliable. Fog is the inverse of comprehension: when the operator cannot read whether the system's horizon-distance estimates are trustworthy, the loop has lost integrity at the comprehension layer.

¹⁷ Forthcoming IH-3 *Confident Misalignment (Horizon-Framed)* (in concurrent drafting; this dispatch). Cites MR-P1[^p1] as canonical; this paper re-describes without redefining.



The four signals correspond, in the IH-4 paper, to the four weather classes (clear, drift, fog, storm) the diagnostic instrument supplies.¹⁸ In this paper the signals are named as the operator-facing surface; in IH-4 they are operationalized as readings the operator continuously consults.

3.5 WHAT THE HORIZON IS NOT

The horizon is **not** a workflow timer. A horizon may correspond to a particular task duration, but the temporal boundary is set by the *operator's pinned commitments*, not by the system's task scheduler. A horizon is **not** a model context window: the substrate boundary may overlap with the context window's edges, but the horizon's substrate dimension covers persistence, tool access, and observability infrastructure as well. A horizon is **not** a session boundary: a single session may span multiple horizons, and a single horizon may extend across multiple sessions. A horizon is **not** an alignment score: scoring a system's output against intent is one operator-side practice, but the horizon is the *envelope* against which scoring is meaningful, not the score itself.

The horizon is, precisely, the named architectural primitive for the temporal-conceptual envelope of operator intent. Naming it is the contribution of this paper; specifying its observable signals is what makes it operationally useful.

4. THE HORIZON WITHIN HGC³AE²'S ROLE ORDERING

The HGC³AE² framework specifies a role ordering: humans govern; context is curated; AI executes; master stamps last.¹⁹ The role ordering applies to the horizon as much as to the intent it bounds.

Humans govern the horizon. The horizon's four boundaries are operator-pinned at governance time: the operator commits the system to operating under a horizon with specific temporal duration, conceptual scope, substrate assumptions, and governance authority. The horizon's *commitment* is an act of governance, not of agentic inference. A system that infers its own horizon is operating outside the role ordering; a system that *reads* a horizon committed by the operator is operating within it.

Curated context maintains horizon situational awareness. The conditions against which the horizon's boundaries are observable — fab calendar, accelerator allocation cycle, model version, persistence-tier health, operator availability, governance-event log — are curated artifacts maintained at Skipjack cadence.²⁰²¹ When the conditions shift, the curated context surfaces the shift; the operator decides whether the shift compresses the horizon or remains within it. The curation discipline is what makes horizon-distance estimation continuous rather than periodic.

Agentic execution operates within the horizon. The system executes against the operator-pinned intent within the operator-pinned horizon. Execution may produce micro-decisions that drift toward a boundary; the architectural commitment is that the system *signals*

18 Forthcoming IH-4 *Epistemic Weather* (in concurrent drafting; this dispatch).

19 Justin H. Kuiper, CISSP, *Mitigating Confident Misalignment in Agentic Systems: The HGC³AE² Framework* (v1.0-preprint, 2026), zenodo.org/records/19869285.

20 Justin H. Kuiper, CISSP, *Mitigating Confident Misalignment in Agentic Systems: The HGC³AE² Framework* (v1.0-preprint, 2026), zenodo.org/records/19869285.

21 Justin H. Kuiper, CISSP, *Agentic Substrate* (v1.0-preprint, 2026), zenodo.org/records/19869291.



the drift rather than crossing the boundary silently. Skipjack's enforcement mechanisms operationalize this commitment (§5).

Runtime verification reads the horizon's state. Verification is not retrospective audit; it is the continuous reading of horizon-state signals while execution proceeds. The Comprehension Instrument's five functions, with the horizon as primitive, become a runtime-verification surface that reads against operator-pinned boundaries. When a function reports a boundary approaching, the architecture fires the appropriate Skipjack enforcement; when a function reports fog, the architecture escalates to the operator before continuing.

The role ordering is not just a static assignment of responsibilities; it is the **closed loop** the architecture commits to. With the horizon as primitive, the loop becomes specifiable: the operator pins, the context curates, the system executes within the pin, the verification layer reads against the pin, and the closure point is the next Skipjack ceremony in which the operator either re-pins or confirms the existing horizon. Without the horizon, the same loop exists but is not specifiable — it operates correctly when conditions cooperate and degrades silently when they do not.

5. THE SKIPJACK PROTOCOL AS HORIZON ENFORCEMENT

The Skipjack Protocol's four enforcement mechanisms — context integrity, structured task decomposition, human-in-the-loop control points, iterative feedback²² — were specified before the horizon was named. Each mechanism becomes operationally precise once the horizon is the primitive it operates against. This section sketches the mapping; the IH-5 capstone develops the full re-description.²³

Context integrity = horizon-substrate gate. The context-integrity mechanism verifies that the context the system is reasoning over has not been corrupted or stale-served. With the horizon as primitive, the mechanism verifies specifically that the *substrate boundary* of the horizon remains valid. A context-integrity failure is, in horizon terms, a substrate-boundary breach.

Structured task decomposition = horizon-distance gate. The structured-decomposition mechanism scopes work to the envelope the operator committed. With the horizon as primitive, the mechanism enforces that each decomposed task remains within the *conceptual boundary* of the horizon. A decomposition that would route work outside the conceptual boundary fails-closed at the gate; the gate is the conceptual-boundary-distance signal made enforceable.

Human-in-the-loop control points = horizon-storm escalation. The human-in-the-loop mechanism inserts the operator at points where the system's self-reports indicate uncertainty above a threshold. With the horizon as primitive, the thresholds are specifiable: storm signals (active boundary crossing) escalate unconditionally; drift signals above a configured rate escalate; fog signals escalate to the operator with metadata about which comprehension function is uncertain.

Iterative feedback = horizon-recalibration loop. The iterative-feedback mechanism uses post-action signals to adjust subsequent action. With the horizon as primitive, the feedback is

22 Justin H. Kuiper, CISSP, *Agile Scrum, Agentics, and the Skipjack Protocol* (v1.0-preprint, 2026), zenodo.org/records/19869293.

23 Forthcoming IH-5 *Skipjack as Horizon Governance* (in concurrent drafting; this dispatch). Cites F2 Skipjack canon and MR-P5[^{p5}] as canonical; this paper re-describes without redefining.



read against horizon boundaries: drift detected in $\$N$ adjusts the horizon-distance estimates for $\$N+1$; storm encountered in $\$N$ triggers governance-time re-pinning before $\$N+1$ fires.

The mapping is not a redefinition. Skipjack's four mechanisms are unchanged in their operational specification. What changes is that each mechanism, with the horizon as primitive, becomes a *gate against a named boundary* rather than a generic alignment safeguard. The architecture acquires a vocabulary in which what was previously implicit can be stated.

6. IMPLICATIONS AND CONCLUSION

The implications of naming the intent horizon extend to operational maturity, tooling investment, and the structure of governed agentic operation as a discipline. Operational maturity acquires a continuous gradation: a system is mature in proportion to how much of its horizon state it can read in real time, and how quickly it can re-pin when boundaries shift. Tooling investment is reframed: dashboards that report system performance against intent become dashboards that report system performance against the *horizon*, with all four boundaries enumerable and inspectable. The discipline is reframed: governed agentic operation is the practice of maintaining horizon-integral execution under shifting boundary conditions.

The core claim: **the intent horizon was always the load-bearing primitive of HGC³AE²; naming it permits structural reasoning about the architectural commitments that previously had to be carried implicitly.** The downstream papers in the *Intent Horizon Theory* series operate on this primitive:

IH-2 Intent Integrity specifies the security model that intent — as the horizon-bounded artifact — requires (CIA + non-repudiation applied to intent itself, not to data).²⁴

IH-3 Confident Misalignment re-describes the dominant failure mode of agentic systems²⁵ through the horizon lens, making the failure mechanically diagnosable.²⁶

IH-4 Epistemic Weather specifies the diagnostic instruments by which operators read horizon state in real time, with four named weather classes.²⁷

IH-5 Skipjack as Horizon Governance is the capstone: Skipjack's four enforcement mechanisms re-described as the operational gates on horizon discipline.²⁸

Returned to HGC³AE² letter by letter through the horizon lens: **H** — humans govern *the horizon*, not just the intent within it. **G** — curated context maintains horizon situational awareness across all four boundaries. **C¹** — cybersecurity protects intent's CIA-NR across horizon boundaries (IH-2's contribution). **C²** — continuity is the property of the horizon, not just the system; horizon recovery after substrate disruption is what continuity operationally requires. **C³** — coherence is horizon-coherence: a unified envelope across substrates, accounts, personas, and workload classes within the operator-pinned horizon. **A** — agentic execution

24 Forthcoming IH-2 *Intent Integrity* (in concurrent drafting; this dispatch).

25 Justin H. Kuiper, CISSP, *Mitigating Confident Misalignment in Agentic Systems: The HGC³AE² Framework* (v1.0-preprint, 2026), zenodo.org/records/19869285.

26 Forthcoming IH-3 *Confident Misalignment (Horizon-Framed)* (in concurrent drafting; this dispatch). Cites MR-P1[^{p1}] as canonical; this paper re-describes without redefining.

27 Forthcoming IH-4 *Epistemic Weather* (in concurrent drafting; this dispatch).

28 Forthcoming IH-5 *Skipjack as Horizon Governance* (in concurrent drafting; this dispatch). Cites F2 Skipjack canon and MR-P5[^{p5}] as canonical; this paper re-describes without redefining.



operates within the horizon, not outside it. **E** — execution is gated against horizon-boundary signals via Skipjack mechanisms (IH-5's contribution). The exponent — the closure — is the next Skipjack ceremony in which the operator re-pins.

This paper closes Paper One of the *Intent Horizon Theory* series. The architecture downstream of this paper is unchanged; the architecture's *diagnostic precision* is changed because the primitive that was always load-bearing is now named.

REFERENCES

— end v0.1-seed draft —

Drafting state: v0.1 IH-1 (chain v0.1 step 1). 6-section canon shape with §2 Problem Set (5 failures) → §3 Definitional Foundation → §4 Role Ordering → §5 Skipjack Mapping → §6 Implications. PR count: 12 (P1, P5, P7, P8, P9, P10 NSQ inheritance + Doctorow Verso · Lakoff/Johnson UChicago · Åström/Hägglund ISA · Mökander *AI and Ethics* + 4 IH-internal forward refs to companion papers). Meets ≥ 12 floor.