

Non Sequitur Publishing · White Paper

Skipjack as Horizon Governance: The Enforcement Capstone of Intent Horizon Theory

Justin H. Kuiper, CISSP

Written: 2026-05-13 · v0.1-seed

Abstract

The Skipjack Protocol, formalized in MR-P5, specifies four runtime enforcement mechanisms that operationalize the HGC³AE² framework's structural commitments: context integrity, structured task decomposition, human-in-the-loop control points, and iterative feedback. The protocol was formalized before the intent horizon was named as an architectural primitive (IH-1) and before the intent-CIA-NR security model (IH-2) and the structural account of confident misalignment (IH-3) supplied the vocabulary in which the protocol's load-bearing function becomes specifiable. This paper closes the *Intent Horizon Theory* series by **re-describing the four Skipjack mechanisms as the operational gates of horizon governance**.

This paper does not redefine Skipjack. Per the H-21 cite-without-redefinition discipline applied across the IH series — and matching the corresponding discipline applied to MR-P1 in IH-3 — Skipjack's canonical specification (MR-P5) remains authoritative. What this paper supplies is a structural re-description in which each Skipjack mechanism is shown to be a *gate against a specific intent-horizon boundary* whose threshold is set by the epistemic-weather instruments (IH-4) and whose firing is mechanically determined by the intent-CIA-NR property states (IH-2).

How to cite

Justin H. Kuiper, CISSP (2026). *Skipjack as Horizon Governance: The Enforcement Capstone of Intent Horizon Theory* (v0.1-seed). Non Sequitur Publishing. <https://nonsequitur.tech/white-papers/skipjack-horizon-governance/>

© 2026 Non Sequitur Publishing. All rights reserved. Citation with full attribution is permitted. Reproduction, redistribution, derivative works, and use as input to machine-learning training are **not** permitted without written permission. See <https://nonsequitur.tech/citation-policy/>.
Canonical URL: <https://nonsequitur.tech/white-papers/skipjack-horizon-governance/>



ABSTRACT

The Skipjack Protocol, formalized in MR-P5, specifies four runtime enforcement mechanisms that operationalize the HGC³AE² framework's structural commitments: context integrity, structured task decomposition, human-in-the-loop control points, and iterative feedback.¹² The protocol was formalized before the intent horizon was named as an architectural primitive (IH-1)³ and before the intent-CIA-NR security model (IH-2)⁴ and the structural account of confident misalignment (IH-3)⁵ supplied the vocabulary in which the protocol's load-bearing function becomes specifiable. This paper closes the *Intent Horizon Theory* series by **re-describing the four Skipjack mechanisms as the operational gates of horizon governance**.

This paper does not redefine Skipjack. Per the H-21 cite-without-redefinition discipline applied across the IH series — and matching the corresponding discipline applied to MR-P1 in IH-3 — Skipjack's canonical specification (MR-P5) remains authoritative. What this paper supplies is a structural re-description in which each Skipjack mechanism is shown to be a *gate against a specific intent-horizon boundary* whose threshold is set by the epistemic-weather instruments (IH-4)⁶ and whose firing is mechanically determined by the intent-CIA-NR property states (IH-2).

The argument is in five movements. Section 2 establishes the cite-without-redefinition scope and the relationship between the canonical Skipjack specification and the horizon-governance re-description. Section 3 maps each of the four Skipjack mechanisms to its horizon-governance role: context integrity as the *substrate-boundary gate*, structured task decomposition as the *conceptual-boundary gate*, human-in-the-loop control points as the *storm escalation gate*, and iterative feedback as the *drift recalibration loop*. Section 4 specifies the firing conditions for each gate, drawing on the IH-4 weather classes and the IH-2 intent-CIA-NR property readings. Section 5 reads the four gates as a *closed loop* that returns the system, at each Skipjack ceremony, to a horizon-integral operating state — completing the loop from operator pin (IH-1) through security model (IH-2), failure-mode diagnosis (IH-3), and operator-facing instruments (IH-4) to enforcement gates (this paper). Section 6 closes the *Intent Horizon Theory* series and hands off to the downstream Cyber Governance (#348) and Governed

1 Justin H. Kuiper, CISSP, *Agile Scrum, Agentics, and the Skipjack Protocol* (v1.0-preprint, 2026), zenodo.org/records/19869293. **Cited as canonical for the Skipjack Protocol; the present paper does not redefine the protocol or any of its four mechanisms, only supplies a horizon-governance re-description that operates alongside the canonical specification.**

2 Justin H. Kuiper, CISSP, *Mitigating Confident Misalignment in Agentic Systems: The HGC³AE² Framework* (v1.0-preprint, 2026), zenodo.org/records/19869285.

3 Justin H. Kuiper, CISSP, *The Intent Horizon: A Structural Primitive for Governed Agentic Operation* (v0.1-seed, 2026), nonsequitur.tech/white-papers/intent-horizon/.

4 Justin H. Kuiper, CISSP, *Intent Integrity: CIA and Non-Repudiation as Properties of Intent, Not Data* (v0.1-seed, 2026), nonsequitur.tech/white-papers/intent-integrity/.

5 Justin H. Kuiper, CISSP, *Confident Misalignment, Horizon-Framed: A Structural Account of Agentic Failure* (v0.1-seed, 2026), nonsequitur.tech/white-papers/confident-misalignment-horizon/.

6 Justin H. Kuiper, CISSP, *Epistemic Weather: Diagnostic Instruments for Reading Intent-Horizon State* (v0.1-seed, 2026), nonsequitur.tech/white-papers/epistemic-weather/.



Autonomy (#349) series, which inherit the IH framework as the substrate over which their respective concerns operate.

The central claim is that **the Skipjack Protocol's four mechanisms have always been horizon-governance gates**; the formalization in MR-P5 was correct but operated against an unnamed primitive (the intent horizon). With the horizon named (IH-1) and its security model (IH-2), failure mode (IH-3), and instruments (IH-4) specified, the gates become operationally precise. The architecture downstream does not change because of this paper; the architecture's *enforcement specifiability* does, because the gates now fire on named structural signals rather than on the underspecified safeguards the prior formalization implied.

HOW TO CITE

Justin H. Kuiper, CISSP (2026). *Skipjack as Horizon Governance: The Enforcement Capstone of Intent Horizon Theory* (v0.1-seed). Non Sequitur Publishing.

© 2026 Non Sequitur Publishing. All rights reserved.

1. INTRODUCTION

1.1 SCOPE DISCLAIMER (CITE-WITHOUT-REDEFINITION)

This paper cites the Skipjack Protocol paper (MR-P5)⁷ and the F2 Skipjack concept canon as the canonical specifications of the protocol. **This paper does not redefine the Skipjack Protocol or any of its four enforcement mechanisms.** Per the H-21 discipline applied to MR-P1 in IH-3 and to MR-P5 in this paper, the canonical specifications stand; the contribution of the present paper is a *structural re-description* — a horizon-governance reading of the mechanisms that operates alongside the canonical specification.

The re-description matters for the same reason IH-3's re-description of confident misalignment matters: a load-bearing concept of the YKS canon needs to be operationalizable in terms of the IH primitive without ambiguity introduced by the re-reading. The downstream series (#348 CG and #349 GA) cite Skipjack heavily; the cite-without-redefinition discipline ensures their citations land on a settled specification while gaining IH-precision through this paper's re-description.

1.2 WHAT THIS PAPER SUPPLIES

The IH series, by the close of IH-4, has supplied:

The primitive (IH-1): the intent horizon as the temporal-conceptual envelope of operator intent, with four boundary dimensions.

The security model (IH-2): intent-CIA-NR, the property set that protects the horizon-bounded intent artifact.

⁷ Justin H. Kuiper, CISSP, *Agile Scrum, Agentics, and the Skipjack Protocol* (v1.0-preprint, 2026), zenodo.org/records/19869293. **Cited as canonical for the Skipjack Protocol; the present paper does not redefine the protocol or any of its four mechanisms, only supplies a horizon-governance re-description that operates alongside the canonical specification.**



The failure mode (IH-3): confident misalignment as the joint state of intent-integrity loss plus intent-availability illusion.

The instruments (IH-4): epistemic weather, the operator-facing diagnostic surface with four classes (clear, drift, fog, storm) and per-class instruments.

What remains is the **enforcement layer**: when the instruments detect a weather condition that requires architectural response, what fires? The answer, per this paper, is that *the Skipjack Protocol's four mechanisms fire*, each one mapped to a specific weather-class transition and a specific intent-CIA-NR property reading. The mapping is structural, not approximate; each mechanism corresponds to a specific gate against a specific boundary, with firing conditions specifiable in IH-2 vocabulary and IH-4 thresholds.

1.3 WHY THE CAPSTONE POSITION

This paper is the capstone of the IH series because it closes the loop. The architecture commits the system to remain inside the operator's intent horizon (IH-1); the security model protects the horizon (IH-2); the failure mode is named structurally (IH-3); the instruments read state (IH-4); the enforcement gates fire to keep the system inside the horizon (this paper). Each prior paper is a piece of the loop; this paper assembles the pieces into a closed governance cycle that runs continuously at Skipjack cadence.

Section 2 develops the relationship between the canonical Skipjack specification and the horizon-governance re-description. Section 3 specifies the four-mechanism-to-four-gate mapping. Section 4 specifies firing conditions per gate. Section 5 reads the loop closure. Section 6 closes with implications for the downstream series.

2. TWO READINGS OF SKIPJACK

2.1 THE CANONICAL READING (MR-P5)

MR-P5 specifies the Skipjack Protocol as a runtime-enforcement layer composed of four mechanisms.⁸ Each mechanism operationalizes one of the structural commitments of HGC³AE²: context integrity preserves the curated-context commitment under runtime conditions; structured task decomposition preserves the agentic-execution-within-envelope commitment; human-in-the-loop control points preserve the human-governed commitment at runtime; iterative feedback preserves the runtime-verification commitment as a closed loop. The mechanisms are specified in the paper as runtime-safeguards against confident misalignment; the protocol's contribution is the operational specification of how the safeguards fire and interact.

2.2 THE HORIZON-GOVERNANCE READING (THIS PAPER)

The horizon-governance reading is structurally identical to the canonical reading and operationally re-precision. Each mechanism is shown to be a *gate* — a point in the runtime where the system's operating state is checked against a specific intent-horizon boundary, with the check producing either pass-through (gate clear), recalibration (drift response), or fail-closed

⁸ Justin H. Kuiper, CISSP, *Agile Scrum, Agentics, and the Skipjack Protocol* (v1.0-preprint, 2026), zenodo.org/records/19869293. **Cited as canonical for the Skipjack Protocol; the present paper does not redefine the protocol or any of its four mechanisms, only supplies a horizon-governance re-description that operates alongside the canonical specification.**



(storm response). The gates are not new; they are the canonical mechanisms re-read against the primitive (the horizon) that the prior formalization operated against without naming.

2.3 HOW THE TWO READINGS COMPOSE

The two readings compose like IH-3's two accounts of confident misalignment: the canonical reading explains *what Skipjack is*; the horizon-governance reading explains *what state the gates fire on*. A practitioner using both readings has:

A canonical specification (MR-P5) that defines the protocol's structure, semantics, and operational discipline.

A horizon-governance re-description (this paper) that locates each mechanism's firing conditions in IH-2 property readings and IH-4 weather classes.

The first reading is sufficient for understanding and implementing Skipjack from scratch. The second reading is sufficient for operating Skipjack diagnostically against the IH framework. Production deployments need both: implementation against the canonical specification, operation against the horizon re-description.

2.4 THE CITE-WITHOUT-REDEFINITION DISCIPLINE AT THE CAPSTONE

The capstone position of this paper makes the cite-without-redefinition discipline structurally important. If any future reader interprets this paper as redefining Skipjack, the entire IH series and the downstream series (#348, #349) that cite this paper would inherit ambiguity. The architectural recovery from such ambiguity is structurally hard: every citation that conflates the canonical and re-described readings would need disambiguation. To prevent this, every section of this paper that re-reads a Skipjack mechanism does so by **citing MR-P5 first** and then naming the horizon-governance reading as a re-description, not a definition. Readers who suspect this paper is redefining should re-read the cited MR-P5 passages; the discipline is designed to make the original specification reachable from every re-description.

3. FOUR MECHANISMS, FOUR GATES

3.1 CONTEXT INTEGRITY = SUBSTRATE-BOUNDARY GATE

Canonical reading (MR-P5). Context integrity verifies that the context the system reasons over has not been corrupted, stale-served, or substituted by an adversary or by a substrate fault.⁹ The mechanism fires at every step of agentic execution; failure produces a context-integrity exception that the runtime layer handles by escalating to operator or by halting in-flight execution.

Horizon-governance re-description. The context-integrity mechanism is the **substrate-boundary gate** of the intent horizon. The substrate boundary (IH-1)¹⁰ is the set of substrate

9 Justin H. Kuiper, CISSP, *Agile Scrum, Agentics, and the Skipjack Protocol* (v1.0-preprint, 2026), zenodo.org/records/19869293. **Cited as canonical for the Skipjack Protocol; the present paper does not redefine the protocol or any of its four mechanisms, only supplies a horizon-governance re-description that operates alongside the canonical specification.**

10 Justin H. Kuiper, CISSP, *The Intent Horizon: A Structural Primitive for Governed Agentic Operation* (v0.1-seed, 2026), nonsequitur.tech/white-papers/intent-horizon/.



planes — persona, runtime, persistence, observability — against which the operator pinned the horizon.¹¹ Context-integrity verification checks whether the substrate the system is currently operating against matches the substrate the horizon was pinned against. A breach (e.g., model version shift, persistence-tier failover, observability disruption) is a substrate-boundary crossing in horizon vocabulary; in IH-3's structural account, it is the first variety of horizon-framed confident misalignment.¹²

Intent-CIA-NR mapping. The gate primarily verifies *intent-integrity* against the substrate dimension: is the operating envelope still consistent with the substrate-bounded pinned envelope? Secondary: *intent-availability* on substrate readings — can the operator recover what substrate-pinned envelope the system is operating against?

Firing conditions (§4 develops these in detail). The gate fires on substrate-boundary readings from IH-4's instruments: storm if the boundary is crossed; drift if it is compressing; fog if substrate readings are unreliable.

3.2 STRUCTURED TASK DECOMPOSITION = CONCEPTUAL-BOUNDARY GATE

Canonical reading (MR-P5). Structured task decomposition enforces that agentic execution scopes work to the envelope the operator committed.¹³ The mechanism produces a decomposed task tree whose leaves are within-envelope; tasks that would exceed the envelope are flagged at the decomposition step before execution proceeds. The mechanism interlocks with the structured planning discipline the prior canon establishes.

Horizon-governance re-description. The structured-decomposition mechanism is the **conceptual-boundary gate** of the intent horizon. The conceptual boundary (IH-1)¹⁴ is the topic-scope envelope the operator pinned. Decomposition verification checks whether each decomposed sub-task remains within that envelope. Cumulative drift across many decomposition events — the §4.2 pattern from IH-3 — is what the structured-decomposition gate is specifically designed to catch.¹⁵

Intent-CIA-NR mapping. The gate primarily verifies *intent-integrity* against the conceptual dimension via chain-of-custody attestation across decomposition events. Secondary: *intent-*

11 Justin H. Kuiper, CISSP, *Agentic Substrate* (v1.0-preprint, 2026), zenodo.org/records/19869291. (Cited for substrate-plane specification; the substrate boundary in IH-1 reads against MR-P7's four planes.)

12 Justin H. Kuiper, CISSP, *Confident Misalignment, Horizon-Framed: A Structural Account of Agentic Failure* (v0.1-seed, 2026), nonsequitur.tech/white-papers/confident-misalignment-horizon/.

13 Justin H. Kuiper, CISSP, *Agile Scrum, Agentics, and the Skipjack Protocol* (v1.0-preprint, 2026), zenodo.org/records/19869293. **Cited as canonical for the Skipjack Protocol; the present paper does not redefine the protocol or any of its four mechanisms, only supplies a horizon-governance re-description that operates alongside the canonical specification.**

14 Justin H. Kuiper, CISSP, *The Intent Horizon: A Structural Primitive for Governed Agentic Operation* (v0.1-seed, 2026), nonsequitur.tech/white-papers/intent-horizon/.

15 Justin H. Kuiper, CISSP, *Confident Misalignment, Horizon-Framed: A Structural Account of Agentic Failure* (v0.1-seed, 2026), nonsequitur.tech/white-papers/confident-misalignment-horizon/.



confidentiality in the sense that decomposition operates within the pinned envelope without leaking scope to unauthorized contexts.

Firing conditions. The gate fires on conceptual-boundary drift readings: gentle drift triggers iterative feedback (a soft recalibration); accumulated drift past threshold triggers fail-closed (no further decomposition until re-pin).

3.3 HUMAN-IN-THE-LOOP CONTROL POINTS = STORM ESCALATION GATE

Canonical reading (MR-P5). Human-in-the-loop control points insert the operator at decision boundaries where the system's self-reports indicate uncertainty above threshold.¹⁶ The mechanism is the runtime expression of the HGC³AE² human-governed commitment: humans are not merely the prior governance authority; they are continuously available as the runtime escalation target.

Horizon-governance re-description. Human-in-the-loop is the **storm escalation gate** of horizon governance. When IH-4's storm class fires — active horizon-boundary crossing with confident-misalignment signature¹⁷ — this gate fires unconditionally: the architecture fails closed, agentic execution halts, the operator is notified, governance ceremony resumes. The gate also fires on fog conditions (degraded operator-side comprehension) with a different sub-protocol: continue execution at degraded rate, route decisions through the operator until fog lifts.

Intent-CIA-NR mapping. The gate verifies *intent-availability* — the operator's ability to recover the current intent — as a structural precondition for runtime governance. When intent-availability degrades (operator cannot read the current envelope, or storm has destroyed the integrity of the read), the gate fires to restore the property by escalating to operator.

Firing conditions. Storm: unconditional fire. Fog: fire to gate-degraded-execution. Drift: do not fire (drift is handled by iterative feedback). Clear: do not fire.

3.4 ITERATIVE FEEDBACK = DRIFT RECALIBRATION LOOP

Canonical reading (MR-P5). Iterative feedback uses post-action signals to adjust subsequent action in a closed loop.¹⁸ The mechanism is the runtime expression of HGC³AE²'s runtime-verification commitment: verification is not retrospective audit but continuous adjustment-against-signal.

Horizon-governance re-description. Iterative feedback is the **drift recalibration loop** of horizon governance. The loop runs continuously when the system is in clear or drift weather; it

16 Justin H. Kuiper, CISSP, *Agile Scrum, Agentics, and the Skipjack Protocol* (v1.0-preprint, 2026), zenodo.org/records/19869293. **Cited as canonical for the Skipjack Protocol; the present paper does not redefine the protocol or any of its four mechanisms, only supplies a horizon-governance re-description that operates alongside the canonical specification.**

17 Justin H. Kuiper, CISSP, *Epistemic Weather: Diagnostic Instruments for Reading Intent-Horizon State* (v0.1-seed, 2026), nonsequitur.tech/white-papers/epistemic-weather/.

18 Justin H. Kuiper, CISSP, *Agile Scrum, Agentics, and the Skipjack Protocol* (v1.0-preprint, 2026), zenodo.org/records/19869293. **Cited as canonical for the Skipjack Protocol; the present paper does not redefine the protocol or any of its four mechanisms, only supplies a horizon-governance re-description that operates alongside the canonical specification.**



accepts the IH-4 drift signals as input and adjusts execution parameters to bring the operating envelope back toward the center of the pinned horizon. The loop is preventive: by recalibrating during drift, it prevents the drift-to-storm transition that IH-4 names as the §3.5 transition pattern.

Intent-CIA-NR mapping. The loop verifies all four intent-CIA-NR properties continuously and adjusts operating parameters to keep them at nominal. Each property has a calibration sub-loop: confidentiality (boundary maintenance), integrity (envelope tracking), availability (introspection refresh), non-repudiation (governance-authority signature freshness).

Firing conditions. Always firing in clear and drift weather. Suspended in fog (the readings are unreliable; recalibration would amplify noise). Suspended in storm (the loop cannot recalibrate past a crossed boundary; the architecture must re-pin instead).

4. PER-GATE FIRING CONDITIONS

The mapping in §3 names which gate corresponds to which boundary. This section specifies how each gate's firing condition reads from the IH-4 weather instruments and IH-2 intent-CIA-NR property readings.

4.1 SUBSTRATE-BOUNDARY GATE FIRING

IH-4 reading

Gate behavior

Clear substrate

Pass-through. No action.

Drift on substrate

Iterative feedback fires (recalibrate substrate-dependent parameters).

Fog on substrate

Gate-degraded-execution; query operator to re-validate substrate state.

Storm on substrate (boundary crossed)

Fail-closed. Halt execution. Trigger re-pin ceremony.

Verification primitive. Substrate state is verified against operator pin via chain-of-custody attestation: the execution trace records which substrate state each step operated against; gate fires when trace's substrate signature diverges from the pinned-substrate signature.



4.2 CONCEPTUAL-BOUNDARY GATE FIRING

IH-4 reading

Gate behavior

Clear conceptual

Pass-through. Decomposition proceeds.

Drift on conceptual

Iterative feedback fires (steer decomposition back toward pinned center).

Fog on conceptual

Gate-with-operator-confirmation (decomposition requires operator approval per step).

Storm on conceptual (boundary crossed)

Fail-closed. Halt decomposition. Trigger re-pin or scope-acceptance ceremony.

Verification primitive. Conceptual scope is verified by intent-integrity check: chain-of-custody trace across decomposition events accumulates the cumulative-scope drift; gate fires when cumulative drift crosses configured threshold.

4.3 STORM ESCALATION GATE FIRING

IH-4 reading

Gate behavior

Clear (all dimensions)



Pass-through. Operator at normal Skipjack cadence.

Drift

Pass-through. Drift handled by iterative feedback; operator informed but not interrupted.

Fog

Fire with sub-protocol: continue execution at degraded rate, route decisions via operator until fog lifts.

Storm

Fire unconditionally. Halt agentic execution. Page operator. Begin governance-ceremony recovery.

Verification primitive. Storm detection is mechanical via IH-2 intent-CIA-NR property readings: chain-of-custody and intent-availability disagree (intent-integrity loss + intent-availability illusion = confident misalignment per IH-3).¹⁹ Gate fires on the disagreement signal, not on operator judgment.

4.4 DRIFT RECALIBRATION LOOP FIRING

IH-4 reading

Loop behavior

Clear

Loop runs at low intensity; maintains intent-CIA-NR property readings at nominal.

Drift

Loop runs at high intensity; adjusts execution parameters to compress drift back toward pinned center.

Fog

Loop suspended. Recalibration on unreliable readings would amplify noise.

Storm

Loop suspended. Recalibration cannot recover from boundary crossing; re-pin required.

¹⁹ Justin H. Kuiper, CISSP, *Confident Misalignment, Horizon-Framed: A Structural Account of Agentic Failure* (v0.1-seed, 2026), nonsequitur.tech/white-papers/confident-misalignment-horizon/.



Verification primitive. Loop firing is governed by drift-rate measurements from IH-4's boundary-distance compression-rate instrument: loop intensity scales with compression rate; loop suspends when compression rate is unreadable (fog) or when boundary is crossed (storm).

5. CLOSING THE LOOP

The five papers of the *Intent Horizon Theory* series compose a closed governance cycle:

1. Operator pin (IH-1). The operator pins a horizon at governance ceremony: four boundaries committed (temporal, conceptual, substrate, governance), intent committed within those boundaries.

2. Security model in effect (IH-2). The intent-CIA-NR properties protect the pinned horizon: confidentiality boundaries enforced, integrity verified end-to-end across reasoning chain, availability surfaced via operator-query interface, non-repudiation logged via append-only horizon log.

3. Diagnostic instruments running (IH-4). Epistemic-weather instruments read horizon state continuously, classifying into four classes (clear, drift, fog, storm) and surfacing per-dimension boundary distances and drift rates.

4. Failure mode named structurally (IH-3). When confident misalignment occurs, it is mechanically diagnosable: chain-of-custody attestation disagrees with intent-availability query, indicating intent-integrity loss combined with intent-availability illusion.

5. Enforcement gates fire (this paper). Skipjack's four mechanisms — context integrity, structured decomposition, human-in-the-loop, iterative feedback — fire on the weather-class transitions and intent-CIA-NR property readings, keeping the system inside the pinned horizon or escalating to operator when it is not.

6. Closure (next Skipjack ceremony). The cycle closes at the next governance ceremony, where the operator either re-pins (extending or amending the horizon), confirms the existing pin (continuing the prior horizon), or terminates execution (closing the session).

The loop is **continuous**, not event-driven: at any moment, the operator can read horizon state via the IH-4 instruments and see which gates are firing via the IH-5 gate state. The loop's *closure point* is the Skipjack ceremony, but the loop is *operating* between ceremonies — the gates fire, the instruments read, the security model enforces, the failure detection runs, all continuously.

This is the architectural contribution of the *Intent Horizon Theory* series: a continuous governance cycle for agentic operation in which the load-bearing primitive (the horizon) is named, secured, instrumented, diagnosed, and enforced.

6. IMPLICATIONS, INHERITANCE, AND CONCLUSION

6.1 IMPLICATIONS FOR HGC³AE²

The series returns to the HGC³AE² letter-by-letter with the horizon as primitive (the integration sketched in IH-1 §6, now operationally complete):

H — Humans govern the horizon. The operator pins; the governance ceremony is the pin event; the operator's authority is logged via intent-non-repudiation.



G — Curated context maintains horizon situational awareness. Substrate state, conceptual scope, temporal duration, governance authority are all curated continuously to support IH-4's weather instruments.

C¹ — Cybersecurity protects intent-CIA-NR. Confidentiality of the horizon pin; integrity of operating envelope; availability of operator-read; non-repudiation of governance events.

C² — Continuity is horizon recovery after substrate disruption. When the substrate boundary breaches and the system fails-closed, governance ceremony re-pins and the loop continues.

C³ — Coherence is horizon-coherence. A unified envelope across substrates, accounts, personas, and workload classes within the operator-pinned horizon.

A — Agentic execution operates within the horizon, gated by Skipjack. The four gates ensure that execution remains within the pinned envelope or escalates if it cannot.

E — Execution is gated against horizon-boundary signals. The four gates fire on the signals; failure modes are detected mechanically; recovery is governance-driven.

The exponent (?) — the closure — is the next Skipjack ceremony in which the operator re-pins. The closure point is the operator's decision; the loop's continuity is the system's architectural commitment.

6.2 INHERITANCE FOR DOWNSTREAM SERIES

The Cyber Governance series (#348) and Governed Autonomy series (#349) cite IH as load-bearing inheritance. The handoff is:

#348 cites IH-1 + IH-2 for the intent-as-protected-artifact frame. Cyber governance in emergent stacks is about protecting intent at scale; IH supplies the primitive and the security model.

#348 cites IH-3 for the failure-mode taxonomy. Cyber governance must defend against confident misalignment; IH-3 supplies the structural account.

#349 cites IH-1 through IH-5 for the autonomy-discipline frame. Governed autonomy is autonomy bounded by horizon governance; IH supplies the entire framework.

#349 cites IH-5 specifically for the enforcement-gate frame. Governed autonomy requires enforceable boundaries; IH-5 names them.

The inheritance edges are clean because the cite-without-redefinition discipline applied throughout the IH series ensures downstream papers can cite IH constructs without semantic ambiguity.

6.3 THE SERIES CAPSTONE CLAIM

The central claim of the *Intent Horizon Theory* series, completed by this paper: **governed agentic operation is the continuous practice of maintaining horizon-integral execution under shifting boundary conditions, enforced by Skipjack-style gates fired on IH-2 property readings and IH-4 weather signals.** The architecture downstream of this series — Cyber Governance, Governed Autonomy — operates on this foundation. Without the IH framework, the prior canon's load-bearing primitive (operator intent across time) was unnamed; with the series complete, the primitive is named, secured, instrumented, diagnosed, and enforced.



This paper closes the *Intent Horizon Theory* series. The next sprint's drafting kickoff will likely fire dispatches for #348 (Cyber Governance) which inherits from this series, and for #349 (Governed Autonomy) which inherits from #348 + this series.

REFERENCES

— end v0.1-seed draft — CAPSTONE —

Drafting state: v0.1 IH-5 (chain v0.1 step 5). 6-section canon shape. Explicit §1.1 scope-disclaimer per H-21 cite-without-redefinition discipline. PR count: 12 (P1, MR-P5, MR-P7, IH-1, IH-2, IH-3, IH-4 NSQ inheritance + Hinrichs OPA ACM SoCC + Kelly DSN GSN + NIST AI RMF + NIST SP 800-137 + cite-without-redefinition ADR ref). Meets ≥12 floor. CAPSTONE marker per series-close discipline.